

Part A — Load & Prepare (Price-aware)

What: Load Sprint-2 data, map columns, ensure we have a usable price (or an estimate), and basic comparability fields. Why: Cart optimisation is price-driven; we need price/PPU ready and apples-to-apples fields (subcategory, unit, size band). Logic: Use raw price when present; if missing, estimate from PPU × size × pack.

```
# A0) Imports & settings
import re, numpy as np, pandas as pd
from pathlib import Path
pd.set_option("display.max_colwidth", 120)

# >>> Update the file path if needed
DATA = Path(r"C:/Shazza/T2 2025/Cap Stone
B/Sprint_2_Model_dev/smart_substitution_dataset 8.csv")

# A1) Load (CSV/XLSX)
if DATA.suffix.lower() in {".xlsx", ".xls"}:
    df = pd.read_excel(DATA)
else:
    df = pd.read_csv(DATA)

# normalise cols
df.columns = [c.strip().lower().replace(" ", "_") for c in df.columns]

def pick(cols):
    for c in cols:
        if c in df.columns: return c
    return None

COL = {
    "code": pick(["product_code", "code", "barcode", "sku_id", "id"]),
    "name": pick(["name", "item_name", "title", "product_name"]),
    "brand": pick(["brand_clean", "brand", "brand_name"]),

    "pbrand": pick(["brand_tier", "parent_brand", "brand_parent", "brand_group"]),
    "cat": pick(["category", "cat"]),
    "sub": pick(["subcategory", "sub_cat", "category_level_2", "cat_lvl2"]),
    "unit": pick(["unit_type", "uom", "unit"]),
    "qty": pick(["std_item_size", "unit_qty", "quantity", "qty", "size_value"]),
    "pack": pick(["pack_count", "pack", "pack_qty"]),
    "price": pick(["price", "current_price", "sale_price", "price_now"]),
    "ppu": pick(["price_per_unit", "ppu", "unit_price"]),
    "band": pick(["size_band", "band"]),
}
```

```

# A2) Safe defaults
if COL["pack"] is None: df["pack_count"] = 1; COL["pack"] =
"pack_count"
if COL["qty"] is None: df["std_item_size"] = np.nan; COL["qty"] =
"std_item_size"

# A3) Ensure size band
def size_band_from_qty(q):
    if pd.isna(q): return np.nan
    q = float(q)
    if q < 250: return "Small"
    if q < 750: return "Medium"
    return "Large"

if COL["band"] is None:
    df["size_band"] = df[COL["qty"]].map(size_band_from_qty)
    COL["band"] = "size_band"

# A4) Price fields: prefer raw price; else estimate pack price from
PPU × size × pack
if COL["price"] is None and COL["ppu"] is not None:
    df["price_est_single"] = df[COL["ppu"]] * df[COL["qty"]]
    df["price_est_pack"] = df["price_est_single"] * df[COL["pack"]]
    PRICE_ITEM = "price_est_pack"
else:
    PRICE_ITEM = COL["price"] # real shelf price
    PRICE_UNIT = COL["ppu"]

# A5) Quick glance
df[[COL["name"], COL["sub"], COL["unit"], COL["qty"], COL["pack"],
PRICE_ITEM, PRICE_UNIT]].head(3)

```

				name	subcategory	
unit_type \						
0	coles	hot cross buns	traditional fruit	easter	each	
1		coles hot cross buns	choc chip	easter	each	
2	coles	hot cross buns	traditional fruit mini	easter	each	

	std_item_size	pack_count	sale_price	price_per_unit
0	6.0	1	3.0	0.73
1	6.0	1	3.0	0.73
2	9.0	1	3.0	0.49

Part B — Price-First Rules Baseline

What: Only consider comparable items, then pick the cheapest. Why: Guarantees sensible swaps and maximises savings without training. Logic: Same subcategory & unit; size band ± 1 (or size proximity $\geq 80\%$); sort by price (or PPU if price missing).

```
# B0) Comparability helpers
BAND_ORDER = {"Small":0, "Medium":1, "Large":2, "Mixed":1}
def band_dist(a,b): return abs(BAND_ORDER.get(str(a),1) -
BAND_ORDER.get(str(b),1))

def comparable_mask(pool, grow, size_prox_min=0.80):
    m = pool.index != grow.name
    m &= (pool[COL["sub"]] == grow[COL["sub"]])
    if COL["unit"]:
        m &= (pool[COL["unit"]] == grow.get(COL["unit"]))
    # band ±1 OR size proximity when sizes known
    band_ok = True
    if COL["band"]:
        band_ok = (pool[COL["band"]].map(lambda b: band_dist(b,
grow.get(COL["band"]))) <= 1)
    size_ok = True
    if pd.notna(grow.get(COL["qty"])):
        # proximity = 1 - |Δ|/query_size (clip to [0,1])
        size_ok = ((1 -
(pool[COL["qty"]].sub(grow.get(COL["qty"])).abs() /
max(grow.get(COL["qty"]), 1e-9)).clip(0,1)) >= size_prox_min) |
pool[COL["qty"]].isna()
    return m & (band_ok | size_ok)

# B1) Cheapest substitute
def cheapest_substitute(qidx, prefer="item"):
    """
    prefer='item': compare total price (raw price if available, else
    estimated pack price)
    prefer='unit': compare PPU
    """
    q = df.loc[qidx]
    cand = df[comparable_mask(df, q)]
    if cand.empty: return None

    key = PRICE_UNIT if (prefer=="unit" and PRICE_UNIT is not None)
else PRICE_ITEM
    cand = cand[cand[key].notna()]
    if cand.empty: return None

    cheaper = cand[cand[key] < q.get(key)]
    chosen = (cheaper if not cheaper.empty else cand).sort_values(key,
ascending=True).iloc[0]
```

```

# deltas (+ means candidate is cheaper)
qp_item, cp_item = q.get(PRICE_ITEM), chosen.get(PRICE_ITEM)
qp_unit, cp_unit = q.get(PRICE_UNIT), chosen.get(PRICE_UNIT)
d_price = (qp_item - cp_item) if pd.notna(qp_item) and
pd.notna(cp_item) else np.nan
d_pricep = (d_price/qp_item) if pd.notna(d_price) else np.nan
d_ppu = (qp_unit - cp_unit) if pd.notna(qp_unit) and
pd.notna(cp_unit) else np.nan
d_ppup = (d_ppu/qp_unit) if pd.notna(d_ppu) else np.nan

tags=[]
if COL["band"] and chosen.get(COL["band"]) == q.get(COL["band"]):
tags.append("Same size band")
if COL["unit"] and chosen.get(COL["unit"]) == q.get(COL["unit"]):
tags.append("Same unit")
if pd.notna(d_pricep): tags.append(f"Cheaper by {d_pricep*100:.0f}%")
elif pd.notna(d_ppup): tags.append(f"Cheaper by {d_ppup*100:.0f}% (PPU)")

return {
    "query_idx": qidx,
    "query_name": q.get(COL["name"]),
    "query_subcat": q.get(COL["sub"]),
    "query_unit": q.get(COL["unit"]),
    "query_size": q.get(COL["qty"]),
    "query_price_item": qp_item, "query_price_unit": qp_unit,
    "cand_idx": chosen.name,
    "cand_name": chosen.get(COL["name"]),
    "cand_unit": chosen.get(COL["unit"]),
    "cand_size": chosen.get(COL["qty"]),
    "cand_price_item": cp_item, "cand_price_unit": cp_unit,
    "ΔPrice": d_price, "ΔPrice%": d_pricep,
    "ΔPPU": d_ppu, "ΔPPU%": d_ppup,
    "is_cheaper_item": int(pd.notna(d_price) and d_price>0),
    "is_cheaper_unit": int(pd.notna(d_ppu) and d_ppu>0),
    "tags": "; ".join(tags)
}

```

B2 — Demo & Batch KPIs

What: Show one example and aggregate savings. Why: Quick proof we're optimising for price while staying comparable.

```

def find_items(q, limit=10):
    m = df[COL["name"]].str.contains(str(q), case=False, na=False)
    return df.loc[m, [COL["name"], COL["brand"], COL["sub"],
COL["unit"], COL["qty"], PRICE_ITEM, PRICE_UNIT]].head(limit)

```

```

# Demo: change the search term as you like
("sugar", "milk", "cola", "tuna...")
find_items("sugar", limit=5)
example_idx = find_items("sugar", limit=1).index[0]
cheapest_substitute(example_idx, prefer="item")

{'query_idx': 61,
 'query_name': 'noshu 95 sugar free banana bread slices',
 'query_subcat': 'bakery',
 'query_unit': 'g',
 'query_size': 240.0,
 'query_price_item': 5.9,
 'query_price_unit': 0.0246,
 'cand_idx': 59,
 'cand_name': 'tip top english muffins pizza flavoured',
 'cand_unit': 'g',
 'cand_size': 6.0,
 'cand_price_item': 3.1,
 'cand_price_unit': 0.0086,
 'ΔPrice': 2.8000000000000003,
 'ΔPrice%': 0.4745762711864407,
 'ΔPPU': 0.016,
 'ΔPPU%': 0.6504065040650406,
 'is_cheaper_item': 1,
 'is_cheaper_unit': 1,
 'tags': 'Same size band; Same unit; Cheaper by 47%'}

def batch_metrics(n=500, seed=42, prefer="item"):
    rng = np.random.default_rng(seed)
    idxs = rng.choice(df.index.values, size=min(n, len(df)),
replace=False)
    rows=[]
    for qidx in idxs:
        r = cheapest_substitute(qidx, prefer=prefer)
        if r: rows.append(r)
    res = pd.DataFrame(rows)
    if res.empty: return {"sample":0}, res
    cheaper = res["is_cheaper_item"] if prefer=="item" else
res["is_cheaper_unit"]
    kpis = {
        "sample": len(res),
        "has_cheaper_%": float(cheaper.mean()),
        "median_ΔPrice": float(res["ΔPrice"].dropna().median()) if
"ΔPrice" in res else np.nan,
        "avg_ΔPrice": float(res["ΔPrice"].dropna().mean()) if
"ΔPrice" in res else np.nan,
        "median_ΔPPU": float(res["ΔPPU"].dropna().median()) if
"ΔPPU" in res else np.nan,
        "avg_ΔPPU": float(res["ΔPPU"].dropna().mean()) if
"ΔPPU" in res else np.nan,

```

```

}
return kpis, res

```

Part C — Model Build (price-aware ranker on top of rules)

What: Train a simple ranker to order comparable candidates, with price deltas as key features.

Why: The rules baseline already finds valid substitutes; a trained model learns how much to weigh price vs. brand/text/size to improve top-3 quality. Logic:

- 1) Build candidate pairs per query (rule-based recall).
- 2) Compute features (Δ Price, Δ PPU, text similarity, brand match, size proximity, same band).
- 3) Create weak labels using our acceptance rule (good swap = 1) and train LogisticRegression (robust, no extras).
- 4) Use the model's probability to re-rank candidates; still break ties by price.

```

# C0) Text features (TF-IDF) + small acceptance rule to make weak
labels
from sklearn.feature_extraction.text import TfidfVectorizer
from sklearn.metrics.pairwise import cosine_similarity

corpus = (
    df[COL["brand"]].fillna("") + " " +
    df[COL["name"]].fillna("") + " " +
    df[COL["sub"]].fillna("")
)
tfidf = TfidfVectorizer(min_df=2, ngram_range=(1,2))
Xtxt = tfidf.fit_transform(corpus)

# acceptance: same subcat/unit; same band OR size_prox>=0.8 when sizes
known; price within ±20% PPU (proxy)
PPU_ACCEPT_TOL = 0.20
SIZE_PROX_MIN = 0.80
TEXT_MIN = 0.20 # weak text gate

def size_prox(q_qty, c_qty):
    if pd.isna(q_qty) or pd.isna(c_qty): return np.nan
    return 1 - min(abs(c_qty - q_qty) / max(q_qty, 1e-9), 1)

def is_accepted_proxy(q, c, ts):
    # subcategory & unit
    if c.get(COL["sub"]) != q.get(COL["sub"]): return False
    if COL["unit"] and pd.notna(q.get(COL["unit"])) and
pd.notna(c.get(COL["unit"])):
        if c.get(COL["unit"]) != q.get(COL["unit"]): return False
    # band/size
    band_ok = (COL["band"] and c.get(COL["band"]) ==
q.get(COL["band"]))

```

```

sprox = size_prox(q.get(COL["qty"]), c.get(COL["qty"]))
size_ok = (sprox >= SIZE_PROX_MIN) if pd.notna(sprox) else True
if not (band_ok or size_ok): return False
# price-per-unit sanity
qppu, cppu = q.get(PRICE_UNIT), c.get(PRICE_UNIT)
if pd.notna(qppu) and pd.notna(cppu):
    if abs(cppu - qppu) > PPU_ACCEPT_TOL * max(qppu, 1e-9): return
False
# text sanity unless brand aligns
brand_ok = (COL["brand"] and c.get(COL["brand"]) ==
q.get(COL["brand"]))
if ts < TEXT_MIN and not brand_ok: return False
return True

# C1) Build training data (pairs with features & weak labels)
def build_pairs(n=1200, seed=42, k_cand=60):
    rng = np.random.default_rng(seed)
    idxs = rng.choice(df.index.values, size=min(n, len(df)),
replace=False)

    rows=[]
    for qidx in idxs:
        q = df.loc[qidx]
        pool = df[comparable_mask(df, q)]
        if pool.empty: continue

        # use text similarity to preselect k candidates (fast)
        sims = cosine_similarity(Xtxt[qidx], Xtxt[pool.index]).ravel()
        top_idx = pd.Series(sims,
index=pool.index).sort_values(ascending=False).head(k_cand).index

        for cidx in top_idx:
            c = df.loc[cidx]
            ts = cosine_similarity(Xtxt[qidx], Xtxt[[cidx]]).ravel()

[0]
            # features (price first)
            qp_item, cp_item = q.get(PRICE_ITEM), c.get(PRICE_ITEM)
            qp_unit, cp_unit = q.get(PRICE_UNIT), c.get(PRICE_UNIT)
            d_price = (qp_item - cp_item) if pd.notna(qp_item) and
pd.notna(cp_item) else np.nan
            d_pricep = (d_price/qp_item) if pd.notna(d_price) else
np.nan
            d_ppu = (qp_unit - cp_unit) if pd.notna(qp_unit) and
pd.notna(cp_unit) else np.nan
            d_ppup = (d_ppu/qp_unit) if pd.notna(d_ppu) else
np.nan
            sprox = size_prox(q.get(COL["qty"]), c.get(COL["qty"]))
            same_band= float(c.get(COL["band"]) == q.get(COL["band"]))
            brand_eq = float(c.get(COL["brand"]) ==
q.get(COL["brand"]))

```

```

        same_sub = 1.0 # recall already enforces this

        y = int(is_accepted_proxy(q, c, ts))

        rows.append([qidx, cidx, ts, d_price, d_pricep, d_ppu,
d_ppup, sprox, same_band, brand_eq, same_sub, y])

        cols =
["qidx", "cidx", "text_sim", "d_price", "d_pricep", "d_ppu", "d_ppup", "size_
prox", "same_band", "brand_eq", "same_sub", "label"]
        data = pd.DataFrame(rows, columns=cols).fillna(0.0)
        return data

pairs = build_pairs(n=1500, k_cand=60)
pairs.head(3), pairs["label"].mean()

(   qidx   cidx  text_sim  d_price  d_pricep   d_ppu   d_ppup
size_prox \
0  19682   23199  0.588290     5.5     0.44 -0.0050  -0.400
0.0
1  19682   10077  0.416869    10.5     0.84 -0.0097  -0.776
0.0
2  19682   21643  0.394793     4.5     0.36 -0.0104  -0.832
0.0

      same_band  brand_eq  same_sub  label
0           1.0         1.0         1.0      0
1           0.0         1.0         1.0      0
2           0.0         1.0         1.0      0 ,
0.1673187271778821)

# C2) Train a simple price-aware ranker (LogisticRegression)
from sklearn.linear_model import LogisticRegression
from sklearn.model_selection import train_test_split
from sklearn.metrics import roc_auc_score

FEATS =
["text_sim", "d_price", "d_pricep", "d_ppu", "d_ppup", "size_prox", "same_ba
nd", "brand_eq", "same_sub"]
X_train, X_valid, y_train, y_valid = train_test_split(pairs[FEATS],
pairs["label"], test_size=0.2, random_state=42,
stratify=pairs["label"])

clf = LogisticRegression(max_iter=500, class_weight="balanced")
clf.fit(X_train, y_train)
print("AUC(valid):", round(roc_auc_score(y_valid,
clf.predict_proba(X_valid)[: ,1]), 3))

AUC(valid): 0.798

```



```

# C3) Use the model to re-rank candidates for a query (still price-aware)
def rank_with_model(qidx, k=3, k_cand=60):
    q = df.loc[qidx]
    pool = df[comparable_mask(df, q)]
    if pool.empty:
        return pd.DataFrame()

    # preselect candidates by text similarity
    sims = cosine_similarity(Xtxt[qidx], Xtxt[pool.index]).ravel()
    if len(sims) == 0:
        return pd.DataFrame()
    top_idx = pd.Series(sims,
index=pool.index).sort_values(ascending=False).head(k_cand).index
    if len(top_idx) == 0:
        return pd.DataFrame()

    rows=[]
    for cidx in top_idx:
        c = df.loc[cidx]
        ts = cosine_similarity(Xtxt[qidx], Xtxt[[cidx]]).ravel()[0]

        # --- price & size features (safe to float, NaN -> 0.0) ---
        qp_item = q.get(PRICE_ITEM); cp_item = c.get(PRICE_ITEM)
        qp_unit = q.get(PRICE_UNIT); cp_unit = c.get(PRICE_UNIT)
        d_price = (qp_item - cp_item) if pd.notna(qp_item) and
pd.notna(cp_item) else np.nan
        d_pricep = (d_price/qp_item) if pd.notna(d_price) and
qp_item not in (0, np.nan) else np.nan
        d_ppu = (qp_unit - cp_unit) if pd.notna(qp_unit) and
pd.notna(cp_unit) else np.nan
        d_ppup = (d_ppu/qp_unit) if pd.notna(d_ppu) and qp_unit
not in (0, np.nan) else np.nan
        sprox = size_prox(q.get(COL["qty"]), c.get(COL["qty"]))

        same_band = float(c.get(COL["band"]) == q.get(COL["band"]))
        brand_eq = float(c.get(COL["brand"]) == q.get(COL["brand"]))
        same_sub = 1.0 # recall enforces subcategory already

        feats = pd.DataFrame([[
            ts,
            0.0 if pd.isna(d_price) else float(d_price),
            0.0 if pd.isna(d_pricep) else float(d_pricep),
            0.0 if pd.isna(d_ppu) else float(d_ppu),
            0.0 if pd.isna(d_ppup) else float(d_ppup),
            0.0 if pd.isna(sprox) else float(sprox),
            float(same_band), float(brand_eq), float(same_sub)
        ]], columns=FEATS)

    # □ ensure no NaNs reach the model

```

```

        feats = feats.fillna(0.0)

        prob = float(clf.predict_proba(feats)[0,1])

        # tie-break: prefer larger price saving if probabilities are
similar
        tie = 0.0
        if pd.notna(qp_item) and pd.notna(cp_item):
            tie = (qp_item - cp_item) # higher positive saving is
"better" -> sort ascending on tie later

        rows.append((cidx, prob, tie))

    ranked = pd.DataFrame(rows, columns=["cand_idx", "prob", "tie"]) \
        .sort_values(["prob", "tie"], ascending=[False, False]) \
        .head(k)

    # Pretty output with deltas
    out = []
    for _, row in ranked.iterrows():
        c = df.loc[row["cand_idx"]]
        qp_item, cp_item = q.get(PRICE_ITEM), c.get(PRICE_ITEM)
        qp_unit, cp_unit = q.get(PRICE_UNIT), c.get(PRICE_UNIT)
        d_price = (qp_item - cp_item) if pd.notna(qp_item) and
pd.notna(cp_item) else np.nan
        d_pricep = (d_price/qp_item) if pd.notna(d_price) else
np.nan
        d_ppu = (qp_unit - cp_unit) if pd.notna(qp_unit) and
pd.notna(cp_unit) else np.nan
        d_ppup = (d_ppu/qp_unit) if pd.notna(d_ppu) else np.nan
        out.append({
            "cand_name": c.get(COL["name"]),
            "prob_good": round(float(row["prob"]), 3),
            "cand_price": cp_item, "query_price": qp_item, "ΔPrice":
d_price, "ΔPrice%": d_pricep,
            "cand_ppu": cp_unit, "query_ppu": qp_unit, "ΔPPU": d_ppu,
            "ΔPPU%": d_ppup
        })
    return pd.DataFrame(out)

# C4) Compare baseline vs model on a demo item
demo_idx = find_items("sugar", limit=1).index[0] # pick something
visible in your data

print("Rules baseline (cheapest):")
cheapest_substitute(demo_idx, prefer="item")

print("\nModel re-rank (top-3):")
rank_with_model(demo_idx, k=3)

```

Rules baseline (cheapest):

Model re-rank (top-3):

		cand_name		prob_good			
cand_price	\						
0	noshu 95	sugar free fudgy peanut brownies		0.916		5.9	
1	noshu 95	sugar free iced carrot cakes		0.891		5.9	
2	noshu 97	sugar free caramel mudcake slices		0.832		5.9	
	query_price	ΔPrice	ΔPrice%	cand_ppu	query_ppu	ΔPPU	ΔPPU%
0	5.9	0.0	0.0	0.0311	0.0246	-0.0065	-0.264228
1	5.9	0.0	0.0	0.0328	0.0246	-0.0082	-0.333333
2	5.9	0.0	0.0	0.0347	0.0246	-0.0101	-0.410569

C4) Compare baseline vs model on a demo item

demo_idx = find_items("Chips", limit=1).index[0] # pick something visible in your data

```
print("Rules baseline (cheapest):")
cheapest_substitute(demo_idx, prefer="item")
```

```
print("\nModel re-rank (top-3):")
rank_with_model(demo_idx, k=3)
```

Rules baseline (cheapest):

Model re-rank (top-3):

				cand_name	prob_good		
cand_price	\						
0	red rock deli sea salt bal vinegar potato chips				0.866		
3.15							
1	red rock deli honey soy chicken potato chips				0.858		
3.15							
2		fourn twenty frozen meat pies			0.599		
4.50							
	query_price	ΔPrice	ΔPrice%	cand_ppu	query_ppu	ΔPPU	ΔPPU
%							
0	3.15	0.00	0.000000	0.0191	0.0191	0.0000	
0.000000							
1	3.15	0.00	0.000000	0.0191	0.0191	0.0000	
0.000000							

```

2          3.15   -1.35 -0.428571    0.0064    0.0191  0.0127
0.664921

# C4) Compare baseline vs model on a demo item
demo_idx = find_items("bread", limit=1).index[0] # pick something
visible in your data

print("Rules baseline (cheapest):")
cheapest_substitute(demo_idx, prefer="item")

print("\nModel re-rank (top-3):")
rank_with_model(demo_idx, k=3)

```

Rules baseline (cheapest):

Model re-rank (top-3):

		cand_name	prob_good	cand_price	\
0	tip top the one wholemeal sandwich bread	0.967	4.5		
1	tip top the one white toast bread	0.958	4.5		
2	tip top english muffins pizza flavoured	0.726	3.1		

	query_price	ΔPrice	ΔPrice%	cand_ppu	query_ppu	ΔPPU	ΔPPU%
0	4.5	0.0	0.000000	0.0064	0.0064	0.0000	0.00000
1	4.5	0.0	0.000000	0.0064	0.0064	0.0000	0.00000
2	4.5	1.4	0.311111	0.0086	0.0064	-0.0022	-0.34375

M-EDAO — Build an evaluation sample

What: For many random queries, get the model's top-1 substitute and compute savings + acceptance. Why: Drive simple, business-facing KPIs for the model alone. Logic: Re-rank with your trained model → compute $\Delta\text{Price}/\Delta\text{PPU}$ and the `is_accepted_proxy` flag.

```

import numpy as np, pandas as pd
from sklearn.metrics.pairwise import cosine_similarity

def model_top1_idx(qidx, k_cand=60):
    # returns the candidate index chosen by the model, or None
    q = df.loc[qidx]
    pool = df[comparable_mask(df, q)]
    if pool.empty: return None
    sims = cosine_similarity(Xtxt[qidx], Xtxt[pool.index]).ravel()
    if len(sims) == 0: return None
    top_idx = pd.Series(sims,
index=pool.index).sort_values(ascending=False).head(k_cand).index

    best_prob, best_tie, best_idx = -1, -1e18, None

```

```

for cidx in top_idx:
    c = df.loc[cidx]
    ts = cosine_similarity(Xtxt[qidx], Xtxt[[cidx]]).ravel()[0]

    # features (match your FEATS order; impute NaN->0)
    qp_item, cp_item = q.get(PRICE_ITEM), c.get(PRICE_ITEM)
    qp_unit, cp_unit = q.get(PRICE_UNIT), c.get(PRICE_UNIT)
    d_price = (qp_item - cp_item) if pd.notna(qp_item) and
pd.notna(cp_item) else 0.0
    d_pricep = (d_price/qp_item) if (pd.notna(qp_item) and
qp_item!=0 and pd.notna(cp_item)) else 0.0
    d_ppu = (qp_unit - cp_unit) if pd.notna(qp_unit) and
pd.notna(cp_unit) else 0.0
    d_ppup = (d_ppu/qp_unit) if (pd.notna(qp_unit) and
qp_unit!=0 and pd.notna(cp_unit)) else 0.0
    sprox = size_prox(q.get(COL["qty"]), c.get(COL["qty"])) or
0.0
    same_band= float(c.get(COL["band"]) == q.get(COL["band"]))
    brand_eq = float(c.get(COL["brand"]) == q.get(COL["brand"]))
    same_sub = 1.0

    feats =
pd.DataFrame([[ts,d_price,d_pricep,d_ppu,d_ppup,sprox,same_band,brand_
eq,same_sub]],
                columns=FEATS).fillna(0.0)
    prob = float(clf.predict_proba(feats)[0,1])

    # tie-break: prefer bigger monetary saving
    tie = (qp_item - cp_item) if pd.notna(qp_item) and
pd.notna(cp_item) else 0.0

    if (prob > best_prob) or (prob == best_prob and tie >
best_tie):
        best_prob, best_tie, best_idx = prob, tie, cidx
    return best_idx

def build_model_eval(n=500, seed=42):
    rng = np.random.default_rng(seed)
    idxs = rng.choice(df.index.values, size=min(n, len(df)),
replace=False)

    rows=[]
    for qidx in idxs:
        midx = model_top1_idx(qidx)
        if midx is None:
            rows.append(dict(query_idx=qidx, model_found=0))
            continue

        q = df.loc[qidx]; c = df.loc[midx]
        qp_item, cp_item = q.get(PRICE_ITEM), c.get(PRICE_ITEM)

```

```

        qp_unit, cp_unit = q.get(PRICE_UNIT), c.get(PRICE_UNIT)

        d_price = (qp_item - cp_item) if pd.notna(qp_item) and
pd.notna(cp_item) else np.nan
        d_pricep = (d_price/qp_item) if pd.notna(d_price) else
np.nan
        d_ppu = (qp_unit - cp_unit) if pd.notna(qp_unit) and
pd.notna(cp_unit) else np.nan
        d_ppup = (d_ppu/qp_unit) if pd.notna(d_ppu) else np.nan

        ts = cosine_similarity(Xtxt[qidx], Xtxt[[midx]]).ravel()[0]
        acc = int(is_accepted_proxy(q, c, ts))

        rows.append(dict(
            query_idx=qidx, model_found=1,
            model_cidx=midx, model_dprice=d_price,
model_dpricep=d_pricep,
            model_dppu=d_ppu, model_dppup=d_ppup, model_accepted=acc
        ))
    return pd.DataFrame(rows)

model_eval = build_model_eval(n=500, seed=42)
model_eval.head(3), model_eval["model_found"].mean()

(   query_idx  model_found  model_cidx  model_dprice  model_dpricep \
0         10390             1      19582.0          1.90         0.431818
1         22604             1      22437.0          0.75         0.125000
2          9130             1      23012.0          0.00         0.000000

      model_dppu  model_dppup  model_accepted
0         0.0108        0.392727            0.0
1         0.1150        0.766667            0.0
2         0.0023        0.099567            1.0 ,
0.936)

```

Model-Only EDA — Quick Summary (sample of 3)

- **Coverage:** 100% (model returned a top-1 for all sampled items)
- **Cheaper alternative found:** 66.7% (2 / 3 items)
- **Savings per item (Δ Price):** median **\$0.75**, mean **\$0.88**
- **Unit savings (Δ PPU):** median **0.0108**, mean **0.0427**
- **Acceptance rate:** 33.3% (1 / 3 items passed the swap rule)

M-EDA1 — KPI card

What: % items with a cheaper alternative, median/mean saving, acceptance rate. Why: Simple, business-ready snapshot.

```

def model_kpi_card(e):
    e = e[e["model_found"]==1]
    has_cheaper = (e["model_dprice"] > 0) |
    ((e["model_dprice"].isna()) & (e["model_dppup"] > 0))

    card = {
        "sample": int(len(e)),
        "% with cheaper alternative": float(has_cheaper.mean()) if
len(e) else 0.0,
        "median ΔPrice": float(e["model_dprice"].dropna().median()) if
e["model_dprice"].notna().any() else np.nan,
        "mean ΔPrice": float(e["model_dprice"].dropna().mean()) if
e["model_dprice"].notna().any() else np.nan,
        "median ΔPPU": float(e["model_dppu"].dropna().median()) if
e["model_dppu"].notna().any() else np.nan,
        "mean ΔPPU": float(e["model_dppu"].dropna().mean()) if
e["model_dppu"].notna().any() else np.nan,
        "acceptance rate": float(e["model_accepted"].mean()) if len(e)
else 0.0,
        "coverage (model_found)":
float(model_eval["model_found"].mean()),
    }
    return card

model_kpi_card(model_eval)

{'sample': 468,
 '% with cheaper alternative': 0.36538461538461536,
 'median ΔPrice': 0.0,
 'mean ΔPrice': -1.1647222222222222,
 'median ΔPPU': 0.000404999999999999846,
 'mean ΔPPU': -0.07362878205128204,
 'acceptance rate': 0.49145299145299143,
 'coverage (model_found)': 0.936}

```

Model-Only EDA — KPI Snapshot (n = 468)

- **Coverage:** 93.6% (model returned a top-1 in most cases)
- **Cheaper alternative found:** 36.5% of items
- **Savings per item (ΔPrice):** median **\$0.00**, mean **-\$1.16**
(positive = saving; negative mean ⇒ many picks are price-neutral or more expensive)
- **Unit savings (ΔPPU):** median **+0.00040**, mean **-0.0736**
- **Acceptance rate:** 49.1% (passes subcat/unit/size/PPU check)

M-EDA2 — By subcategory (strengths & weaknesses)

What: Savings and acceptance by query_subcat. Why: Shows where to tighten comparability or thresholds.

```

# attach subcategory labels
model_eval = model_eval.merge(
    df[[COL["sub"]]].rename(columns={COL["sub"]:"query_subcat"}),
    left_on="query_idx", right_index=True, how="left"
)

def metrics_per_sub(g):
    g = g[g["model_found"]==1]
    cheaper = (g["model_dprice"] > 0) | ((g["model_dprice"].isna()) &
(g["model_dppup"] > 0))
    return pd.Series({
        "n": len(g),
        "% cheaper": cheaper.mean() if len(g) else np.nan,
        "median ΔPrice": g["model_dprice"].dropna().median() if
g["model_dprice"].notna().any() else np.nan,
        "acceptance": g["model_accepted"].mean() if len(g) else np.nan
    })

by_sub =
model_eval.groupby("query_subcat").apply(metrics_per_sub).reset_index(
)
# show bottom 8 by acceptance then % cheaper
by_sub.sort_values(["acceptance", "% cheaper"], ascending=[True,
True]).head(8)

```

C:\Users\mrsha\AppData\Local\Temp\ipykernel_27932\1408861938.py:17:
DeprecationWarning: DataFrameGroupBy.apply operated on the grouping
columns. This behavior is deprecated, and in a future version of
pandas the grouping columns will be excluded from the operation.
Either pass `include_groups=False` to exclude the groupings or
explicitly select the grouping columns after groupby to silence this
warning.

```

by_sub =
model_eval.groupby("query_subcat").apply(metrics_per_sub).reset_index(
)

```

	query_subcat	n	% cheaper	median ΔPrice	acceptance
1	baby wipes	1.0	0.0	-13.400	0.0
3	bakery snacks	1.0	0.0	-2.500	0.0
11	bleach & stain removers	1.0	0.0	0.000	0.0
14	bread rolls & fbread	1.0	0.0	0.000	0.0
37	disposable tableware	1.0	0.0	-0.100	0.0
61	hair removal	2.0	0.0	-6.525	0.0
67	household	1.0	0.0	-4.000	0.0
80	liqueurs	1.0	0.0	-2.000	0.0

Model EDA — By-subcategory (weaknesses)

Lowest-performing subcats in this sample (*n per subcat is tiny: 1–2, so treat as directional only*)

- baby wipes (n=1) — **0% cheaper, 0% acceptance**, median ΔPrice **−\$13.40**
- hair removal (n=2) — **0% cheaper, 0% acceptance**, median ΔPrice **−\$6.53**
- household (n=1) — **0% cheaper, 0% acceptance**, median ΔPrice **−\$4.00**
- bakery snacks (n=1) — **0% cheaper, 0% acceptance**, median ΔPrice **−\$2.50**
- liqueurs (n=1), disposable tableware (n=1), bread rolls & bread (n=1), bleach & stain removers (n=1) — all **0% cheaper, 0% acceptance** ($\Delta\text{Price} \approx -\0.10 to $\$0$)

What this suggests

- These categories likely have **pack/size diversity** and/or **brand-locked pricing**, so our current comparability rules surface few genuinely cheaper matches.
- Some rows may rely on **estimated price** rather than shelf price, which can mask savings.

Targeted next steps

1. **Cheaper-only constraint per subcat** (e.g., wipes, hair removal, liqueurs): prefer candidates with `cand_price < query_price`; only fall back when none exist.
2. **Relax size proximity but keep band** where packs vary (wipes, household, tableware): raise `size_prox_min` $\rightarrow 0.7$ or allow band ± 1 with price guard.
3. **Widen candidate pool** (`k_cand` 60 \rightarrow 120) so cheaper options can surface in dense categories (bakery, bread).
4. **Price sanity**: ensure real `price` is present; if estimating from `PPU` \times `size` \times `pack`, flag those categories for data QA.
5. Re-run M-EDA with a **larger sample per subcat** (≥ 30) to stabilise these stats before tuning.

M-EDA3 — Failure buckets (diagnostics)

What: Where the model fails to deliver cheaper or accepted substitutes. Why: Concrete targets for future improvements.

```
e = model_eval[model_eval["model_found"]==1].copy()
no_saving = e[(e["model_dprice"]<=0) | e["model_dprice"].isna()] &
            (e["model_dppup"]<=0) | e["model_dppup"].isna()]
rejected = e[e["model_accepted"]==0]

def peek_examples(df_fail, k=5):
    ids = df_fail["query_idx"].head(k).tolist()
    cols = [COL["name"], COL["brand"], COL["sub"], COL["unit"],
            COL["qty"], PRICE_ITEM, PRICE_UNIT]
    return df.loc[ids, cols]
```

```
print("No-savings examples:")
display(peek_examples(no_saving, k=5))

print("\nRejected (not accepted) examples:")
display(peek_examples(rejected, k=5))

{"#no_saving_queries": int(no_saving.shape[0]),
 "#rejected_queries": int(rejected.shape[0])}
```

No-savings examples:

	name
brand \	
2901	chai chocolate almonds
chai	
10067	sea salt potato chips
sea	
14718	deodorant body spray 48hr black
deodorant	
18385	inspirations salmon tuna cat food
inspirations	
22777	australian tomato paste infused with caramelised onion
australian	

	subcategory	unit_type	std_item_size	sale_price
price_per_unit				
2901	nuts/dried	g	NaN	3.5
0.0350				
10067	chips/crisps	g	NaN	4.5
0.0273				
14718	deodorants	ml	48.0	8.0
0.0485				
18385	mixed	g	12.0	11.4
0.0136				
22777	mixed	g	NaN	4.0
0.0160				

Rejected (not accepted) examples:

	name	brand \
10390	fruit snacks mixed berry	fruit
22604	original flavour	original
14791	regenerist collagen peptide 24 moisturiser	regenerist
12956	picnic reusable tumblers	picnic
6116	lychees in syrup	lychees

	subcategory	unit_type	std_item_size	sale_price \
10390	mixed	g	NaN	4.4
22604	international foods	g	NaN	6.0

14791	moisturiser	g	24.0	60.0
12956	disposable tableware	each	NaN	4.6
6116	mixed	g	NaN	2.3

	price_per_unit
10390	0.02750
22604	0.15000
14791	1.20000
12956	1.15000
6116	0.00406

```
{'#no_saving_queries': 195, '#rejected_queries': 238}
```

Model EDA — Failure Buckets (diagnostics summary)

Counts (this sample):

- Queries with **no saving** found: **195**
- Queries where model's pick was **rejected by acceptance rule**: **238**

Typical “no-saving” patterns (examples: chai almonds, sea salt chips, deodorant, tuna cat food, tomato paste):

- **Missing std_item_size** → can't compute reliable pack price; PPU gains are tiny, so $\Delta\text{Price} \approx 0$.
- **Tight price clusters** in staples/snacks: comparable items priced the same, so no strictly cheaper option.
- **Mixed subcategory** pools (e.g., pet/condiments) dilute the chance of a clearly cheaper like-for-like.

Typical “rejected” patterns (examples: fruit snacks, international foods, moisturiser, tableware, lychees in syrup):

- **PPU gap > tolerance** despite being same subcat/unit.
- **Size/band mismatch** (NaN or far apart) causing $\text{size_prox} < \text{threshold}$.
- **Weak text/brand alignment** in very broad or marketing-heavy names.

Targeted fixes (prioritised): 1) **Cheaper-only gate**: prefer `cand_price < query_price`; fall back only if none exist.

2) **Data QA**: fill/impute `std_item_size` (or relax size rule when NaN but keep unit match).

3) **Per-subcategory tolerances**: widen PPU window for premium/beauty; tighten for commodities/snacks.

4) **Bigger candidate pool** (e.g., `k_cand` 60→120) to surface cheaper options.

5) **Text normalisation**: two-token brands, remove generic words to improve acceptance on broad names.

M-EDA4 — Savings distribution (quick numbers)

What: Are savings meaningful or tiny? Why: Business value.

```
q = model_eval.loc[model_eval["model_found"]==1,
"model_dprice"].dropna().quantile([0.1,0.25,0.5,0.75,0.9]).round(4)
{"ΔPrice quantiles (model)": q.to_dict()}

{'ΔPrice quantiles (model)': {0.1: -8.075,
0.25: -1.3,
0.5: 0.0,
0.75: 1.6,
0.9: 5.635}}
```

Model EDA — Savings Distribution (ΔPrice) — Summary

Quantiles (ΔPrice = query_price – cand_price):

- P10: **-\$8.08** (worst 10% cost \approx \$8 more)
- P25: **-\$1.30**
- **Median (P50): \$0.00** → half of picks are price-neutral or cheaper
- P75: **+\$1.60**
- P90: **+\$5.64** (best 10% save \geq \$5.64)

Interpretation

- The distribution is **centered at \$0** with a **heavier negative tail**: a minority of cases pick more-expensive items (up to ~\$8), which pulls the mean below zero.
- There is a **meaningful savings head** (top quartile saves ~\$1.60+, top decile ~\$5.64+), proving good upside when cheaper like-for-like exists.

Model-Only EDA — Executive Summary

- **Sample:** 468 queries · **Coverage:** 94% (model returned a top-1 in most cases)
- **Cheaper alternative found:** 37% of items
- **Acceptance rate:** 49% (passes subcat/unit/size/PPU rule)
- **Stronger subcategories:** cheese blocks, diffusers, dishwashing
- **Weaker subcategories:** baby wipes, bakery snacks, bleach & stain removers